Stat243: Final group project, Due Friday Dec. 13, 1 pm

December 9, 2013

The project will be done in groups of 3-4, with students assigned randomly but balanced between statistics and non-statistics students.

A few comments. First, the project, when split amongst the group members, is not intended to be a huge undertaking. The goals of the project are to give you experience in working collaboratively and developing a well-designed, well-tested piece of software.

The project will be graded as a letter grade and will count for about as much as a problem set in your final grade.

Formatting requirements

Your solution to the problem should have two parts:

1. The code of your solution, including (a) a primary function that carries out the simulation, (b) function that carries out the formal tests you applied to your function, and (c) auxiliary functions used by the primary function and testing function. All code should be submitted to bSpace as a single text file with comments clearly delineating different pieces of the code. Please call the file "ars.R". Please call your main function "ars" and your test function "test". Of course in your development you could use many files, but please combine for the final one.

2. A paper document describing your solution, prepared in LaTeX or R Markdown. The description does not need to be more than a couple pages, but should describe the approach you took in terms of functions/modularity/object-oriented programming, and the testing that you carried out. It must include a paragraph describing the specific contributions of each team member and which person/people were responsible for each component of the work. Please submit the paper document to me - either directly to me, under my door, or in my mailbox. The deadline is firm because I need to grade them Friday afternoon and during the day Saturday. On your paper solution, please indicate the last name of the group member who submitted the R code so that I know where to look on bSpace. Make sure to include the requested paragraph describing which project member did what.

3. For the help information on your main function, this can either be part of your R code as indicated below or if a PDF, can be stapled to your paper solution.

For a given subtask of the problem, if you find good code available, you may use it as a modular component of your code provided it does not constitute too large a part of your solution and provided you test the code. Consult me with questions on this matter.

You should use Git to manage your collaboration. You could use Github or another resource online or see statistics.berkeley.edu/computing/git for how to do this on the SCF.

You should start the process by mapping out the modular components you need to write and how they will fit together, as well as what the primary function will do. After one person writes a component, another

person on the team should test it and, with the original coder, improve it. You might also consider an approach called pair programming (see the Wikipedia entry) where two people write code side by side.

Problem

Your task is to implement an adaptive-rejection sampler, described in class in Unit 12 and with details in Gilks et al. (1992) - the PDF is in the project directory on github. I'd suggest you implement the tangent approach rather than the secant approach, but the latter is fine too. Your work should include the following components:

1. Your solution should allow the user to provide reasonable inputs, including the number of points to sample, and should check the inputs for validity. The primary input should be an R function or R expression that calculates the (possibly unnormalized) density of the distribution of interest in a vectorized fashion (e.g., many of the "d" functions in R, such as "dnorm" are legitimate inputs). Your code should include numerical checks that catch cases of non-log-concave densities as the calculations proceed. (I.e., you do not need to come up with an overall test of the input density, but you should be able to do some checks that will catch cases where the upper and lower bounds are not actually bounding the density.)

2. Formal testing is required (see the previous section), with a set of tests where results are compared to some known truth. You should have tests for the overall function, and for any modules that do anything complicated. Given the output is stochastic, how to do this will require some thought. The output of your testing function should be clear and interpretable. I.e., when I run your test function, it should print informative messages of what it is doing and whether the test was passed or failed.

3. Your solution should involve modular code, with functions or OOP methods that implement discrete tasks. You should have an overall design and style that is consistent across the components, in terms of functions vs. OOP methods, naming of objects, etc.

4. In terms of efficiency, the algorithm is inherently sequential. However, you should try to vectorize as much as possible. One possibility in terms of the overall calculation is that you could generate a vector of samples based on the upper envelope. Then determine where the points that require evaluation of $f(x)$. All points up to the first of those points can be generated before changing the envelope, at which point you would throw away the remaining points. How many points you generate at once might vary depending on how far along you are in generating the number of points requested by the user. Or you may think of other tricks.

5. Your solution should include help/manual information as follows: (a) basic doc strings for any function you create (i.e., include comments at the beginning of the function), as well as commenting of code as appropriate. (b) An overall help page for the primary function. This could be in the form of comments in the R code in the form used in the *roxygen2* package (see http://adv-r.had.co.nz/Documenting-functions.html for an example) or simply a PDF (e.g., see the help info on, say, p. 29 of http://www.cran.r-project.org/web/packages/spam/spam.pdf for help on the *chol()* function in the *spam* package). I'd prefer the first (and that is likely to be easier), but the latter is fine too.